

Building intuition for performance

Roel Janssen

November 17, 2023



Wait a second

How much can a computer do in a second?

<https://computers-are-fast.github.io/>

Some numbers

execute typical instruction	1/1,000,000,000 sec = 1 nanosec
fetch from L1 cache memory	0.5 nanosec
branch misprediction	5 nanosec
fetch from L2 cache memory	7 nanosec
Mutex lock/unlock	25 nanosec
fetch from main memory	100 nanosec
send 2K bytes over 1Gbps network	20,000 nanosec
read 1MB sequentially from memory	250,000 nanosec
fetch from new disk location (seek)	8,000,000 nanosec
read 1MB sequentially from disk	20,000,000 nanosec
send packet US to Europe and back	150 milliseconds = 150,000,000 nanosec

Big-O notation




Figure: Image source:

medium.com/dataseries/how-to-calculate-time-complexity-with-big-o-notation-9afe33aa4c46

Intuition test

```
simple_call_1 <- function (number) { return (number * number) }  
sapply(numbers_1, simple_call_1)  
  
simple_call_2 <- function (numbers) {  
  for (index in 1:length(numbers)) {  
    numbers[index] <- numbers[index] * numbers[index]  
  }  
  return (numbers)  
}
```

Intuition test

```
numbers_1 <- runif(1000000) * 100  
numbers_2 <- numbers_1  
  
object.size(numbers_2)  
# => 8000048 bytes
```

Execution time and space (1/2)

```
library(bench)
res_sapply <- bench::mark(numbers_3 <- sapply(numbers_1, simple_call_1))
res_for <- bench::mark(numbers_4 <- simple_call_2 (numbers_2))

res_sapply
res_for
```

Execution time and space (2/2)

	memory allocated	median time spent
simple_call_1	30.9MB	2.37s
simple_call_2	7.65MB	136ms

Function call profiling (1/2)

```
Rprof("simple_call_1.out")
numbers_3 <- sapply(numbers_1, simple_call_1)
Rprof("simple_call_2.out")
numbers_4 <- simple_call_2 (numbers_2)
Rprof(NULL)

head(summaryRprof("simple_call_1.out")[[ "by.total" ]])
head(summaryRprof("simple_call_2.out")[[ "by.total" ]])
```

Function call profiling (2/2)

```
> head(summaryRprof("simple_call_1.out"))[["by.total"]]
      total.time total.pct self.time self.pct
"sapply"         1.86     98.94      0.00      0.00
"lapply"         1.72     91.49      1.36     72.34
"FUN"            0.36     19.15      0.36     19.15
"simplify2array" 0.14      7.45      0.00      0.00
"unlist"         0.08      4.26      0.08      4.26
"unique"         0.06      3.19      0.00      0.00
> head(summaryRprof("simple_call_2.out"))[["by.total"]]
      total.time total.pct self.time self.pct
"simple_call_2"  0.12      100      0.12     100
>
```

Memory allocation profiling (1/3)

```
library(profmem)
capabilities("profmem") # Check if your R can do memory profiling

options(profmem.threshold = 0)
mem_details_sapply <- profmem({ numbers_3 <- sapply(numbers_1, simple_call_1) })
mem_details_for    <- profmem({ numbers_4 <- simple_call_2 (numbers_2) })

mem_details_sapply
mem_details_for
```

Memory allocation profiling (2/3)

```
> mem_details_sapply
```

```
...
```

```
Memory allocations:
```

```
Number of 'new page' entries not displayed: 9098
```

	what	bytes	calls
1	alloc	8000048	sapply() -> lapply()
9100	alloc	4000048	sapply() -> simplify2array() -> unique() -> lengths()
9101	alloc	8388656	sapply() -> simplify2array() -> unique() -> unique.default()
9102	alloc	4000048	sapply() -> simplify2array() -> unique() -> unique.default()
9103	alloc	8000048	sapply() -> simplify2array() -> unlist()
total		32388848	

Memory allocation profiling (3/3)

```
> mem_details_for
...
Memory allocations:
      what    bytes      calls
1    alloc 8000048 simple_call_2()
total      8000048
> object.size(numbers_3)
8000048 bytes
>
```

More reading

Hadley Wickham on performance in R

<http://adv-r.had.co.nz/Performance.html>